

Derived Datatypes

Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

<http://creativecommons.org/licenses/by-nc-sa/4.0/>

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Acknowledge EPCC as follows: “© EPCC, The University of Edinburgh, www.epcc.ed.ac.uk”

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.

MPI Datatypes

- Basic types
- Derived types
 - vectors
 - structs
 - others

Basic datatypes


```
int x[10];  
INTEGER:: x(10);  
  
// send all 10 values  
MPI_Send(x, 10, MPI_INT, ...);  
MPI_SEND(x, 10, MPI_INTEGER, ...)
```



```
// send first 4 values  
MPI_Send(&x[0], 4, ...);  
MPI_SEND(x(1), 4, ...)
```



```
// send 5th, 6th, 7th, 8th  
MPI_Send(&x[4], 4, ...);  
MPI_SEND(x(5), 4, ...)
```



```
// ??  
struct mystruct x[10];  
type(mytype) :: x(10)
```



Motivation

- Send / Recv calls need a datatype argument
 - pre-defined values exist for pre-defined language types
 - e.g. `double <-> MPI_DOUBLE; int <-> MPI_INT`
- What about types defined by a program?
 - e.g. structures (in C) or user-defined types (Fortran)
- **Send / Recv** calls take a count parameter
 - what about data that isn't contiguous in memory?
 - e.g. subsections of 2D arrays

Approach

- Can define new types in MPI
 - user calls setup routines to describe new data type to MPI
 - remember, MPI is a library and NOT a compiler!
 - MPI returns a new data type handle
 - store this value in a variable, e.g. **MPI_MY_NEWTYPE**
- Derived types have same status as pre-defined
 - can use in any message-passing call
- Some care needed for reduction operations
 - user must also define a new **MPI_Op** appropriate to the new data type to tell MPI how to combine them

Defining types

- All derived types stored by MPI as a list of basic types and displacements (in bytes)
 - for a structure, types may be different
 - for an array subsection, types will be the same
- User can define new derived types in terms of both basic types and other derived types

Derived Data types - Type

basic datatype 0	displacement of datatype 0
basic datatype 1	displacement of datatype 1
...	...
basic datatype n-1	displacement of datatype n-1

Contiguous Data

- The simplest derived datatype consists of a number of contiguous items of the same datatype.

- C:

```
int MPI_Type_contiguous( int count,  
                        MPI_Datatype oldtype,  
                        MPI_Datatype *newtype)
```

- Fortran:

```
MPI_TYPE_CONTIGUOUS (COUNT, OLDTYPE,  
                    NEWTYPE, IERROR)
```

```
INTEGER COUNT, OLDTYPE, NEWTYPE, IERROR
```

Use of contiguous

- May make program clearer to read
- Imagine sending a block of 4 integers
 - use `MPI_Ssend` with `MPI_INT` / `MPI_INTEGER` and `count = 4`
- Or ...
 - define a new contiguous type of 4 integers called `BLOCK4`
 - use `MPI_Ssend` with `type=BLOCK4` and `count = 1`
- May also be useful intermediate stage in building more complicated types
 - i.e. later used in definition of another derived type

Committing a datatype

- Once a datatype has been constructed, it needs to be committed before it is used in a message-passing call
- This is done using **MPI_TYPE_COMMIT**

- C:

```
int MPI_Type_commit (MPI_Datatype *datatype)
```

- Fortran:

```
MPI_TYPE_COMMIT (DATATYPE, IERROR)  
INTEGER DATATYPE, IERROR
```

Example

```
MPI_Datatype BLOCK4;  
int x[8];
```

```
MPI_Type_contiguous(4, MPI_INT, &BLOCK4);  
MPI_Type_commit(&BLOCK4);
```

```
// Send all 8 integers using the new type  
MPI_Ssend(x, 2, BLOCK4, dest, tag, comm);
```

```
// This call sends the same data  
MPI_Ssend(x, 8, MPI_INT, dest, tag, comm);
```


Vector Datatype Example

Oldtype



5 element stride
between blocks

Newtype



- $\text{count} = 2$
- $\text{stride} = 5$
- $\text{blocklength} = 3$

What is a vector type?

- Why is a pattern with blocks and gaps useful?

A vector type corresponds to a subsection of a 2D array

- Think about how arrays are stored in memory
 - unfortunately, different conventions for C and Fortran!
 - must use statically allocated arrays in C because dynamically allocated arrays (using `malloc`) have no defined storage format
 - In Fortran, can use either static or allocatable arrays

Coordinate System (how I draw arrays)

$x[i][j]$

j



i

$x(i, j)$

$x[0][3]$	$x[1][3]$	$x[2][3]$	$x[3][3]$
$x[0][2]$	$x[1][2]$	$x[2][2]$	$x[3][2]$
$x[0][1]$	$x[1][1]$	$x[2][1]$	$x[3][1]$
$x[0][0]$	$x[1][0]$	$x[2][0]$	$x[3][0]$

$x(1, 4)$	$x(2, 4)$	$x(3, 4)$	$x(4, 4)$
$x(1, 3)$	$x(2, 3)$	$x(3, 3)$	$x(4, 3)$
$x(1, 2)$	$x(2, 2)$	$x(3, 2)$	$x(4, 2)$
$x(1, 1)$	$x(2, 1)$	$x(3, 1)$	$x(4, 1)$

Memory Layout

- You can choose to draw arrays however you like – how you draw them does not change reality!

First index i	Second index j	Format
right	up	coordinates
down	right	matrix
right	down	graphics (scan lines)

- Regardless of how you draw them, the layout in memory is:
 - $x[i][j]$ is followed by $x[i][j+1]$ (in C)
 - $x(i, j)$ is followed by $x(i+1, j)$ (in Fortran)
 - if you create arrays with malloc in C/C++ things are more complicated ...
- Depending on how you draw them, this can appear “row major” or “column major”

Array Layout in Memory

C: $\mathbf{x}[16]$

F: $\mathbf{x}(16)$

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----	----

C: $\mathbf{x}[4][4]$

F: $\mathbf{x}(4, 4)$

4	8	12	16
3	7	11	15
2	6	10	14
1	5	9	13




13	14	15	16
9	10	11	12
5	6	7	8
1	2	3	4

- Data is contiguous in memory
 - different conventions for mapping 2D to 1D arrays in C and Fortran

C example


C: **x**[5][4]



- A 3 x 2 subsection of a 5 x 4 array
 - three blocks of two elements separated by gaps of two



Fortran example

F: $\mathbf{x}(5, 4)$



- A 3 x 2 subsection of a 5 x 4 array
 - two blocks of three elements separated by gaps of two

Equivalent Vector Datatypes



Constructing a Vector Datatype

- C:

```
int MPI_Type_vector (int count,  
                    int blocklength, int stride,  
                    MPI_Datatype oldtype,  
                    MPI_Datatype *newtype)
```

- Fortran:

```
MPI_TYPE_VECTOR (COUNT, BLOCKLENGTH,  
                STRIDE, OLDTYPE, NEWTYPE, IERROR)
```

Sending a vector

- Have defined a **3x2** subsection of a **5x4** array
 - but not defined WHICH subsection
 - is it the bottom left-hand corner? top-right?
- Data that is sent depends on what buffer you pass to the send routines
 - pass the address of the first element that should be sent



Vectors in send routines

```
MPI_Ssend(&x[1][1], 1, vector3x2, ...);
```

```
MPI_SSEND(x(2,2) , 1, vector3x2, ...)
```

```
MPI_Ssend(&x[2][1], 1, vector3x2, ...);
```

```
MPI_SSEND(x(3,2) , 1, vector3x2, ...)
```



Structures

- Can define compound objects in C and Fortran

<pre>struct compound { int ival; double dval[3]; };</pre>	<pre>type compound integer :: ival double precision :: dval(3) end type compound</pre>
--	---

- Storage format NOT defined by the language
 - different compilers do different things
 - e.g. insert arbitrary padding between successive elements
 - need to tell MPI the byte displacements of every element

Constructing a Struct Datatype

- C:

```
int MPI_Type_create_struct (int count,  
    int *array_of_blocklengths,  
    MPI_Aint *array_of_displacements,  
    MPI_Datatype *array_of_types,  
    MPI_Datatype *newtype)
```

- Fortran:

```
MPI_TYPE_CREATE_STRUCT (COUNT,  
    ARRAY_OF_BLOCKLENGTHS,  
    ARRAY_OF_DISPLACEMENTS,  
    ARRAY_OF_TYPES, NEWTYPE, IERROR)
```

Struct Datatype Example

- `count = 2`
- `array_of_blocklengths[0] = 1`
- `array_of_types[0] = MPI_INT`
- `array_of_blocklengths[1] = 3`
- `array_of_types[1] = MPI_DOUBLE`
- But how do we compute the displacements?
 - need to create a compound variable in our program
 - explicitly compute memory addresses of every member
 - subtract addresses to get displacements from origin

Address of a Variable

- C:

```
int MPI_Get_address (void *location,  
                    MPI_Aint *address) ;
```

- Fortran:

```
MPI_GET_ADDRESS (LOCATION, ADDRESS, IERROR)
```

```
<type> LOCATION (*)
```

```
INTEGER (KIND=MPI_ADDRESS_KIND) ADDRESS
```

```
INTEGER IERROR
```

Example

```
struct compound x, y;  
MPI_Datatype MPI_COMPOUND;  
MPI_Aint array_of_displacements[2];  
MPI_Aint addr, addr2;  
  
MPI_Get_address(&x.ival, &addr1);  
MPI_Get_address(&x.dval, &addr2);  
array_of_displacements[0] = 0;  
array_of_displacements[1] = addr2 - addr1;  
  
MPI_Type_create_struct (count, array_of_blocklengths,  
    array_of_displacements, array_of_types, &MPI_COMPOUND);  
MPI_Type_commit(&MPI_COMPOUND);  
  
MPI_Ssend(&y, 1, MPI_COMPOUND, dest, tag, comm);
```

Exercise

Derived Datatypes

- See Exercise 8 on the sheet
- Modify the passing-around-a-ring exercise.
- Calculate two separate sums:
 - rank integer sum, as before
 - rank floating point sum
- Use a **struct datatype** for this.
- If you are a Fortran programmer unfamiliar with Fortran derived types then jump to exercise 8.2
 - illustrates the use of **MPI_Type_vector**